

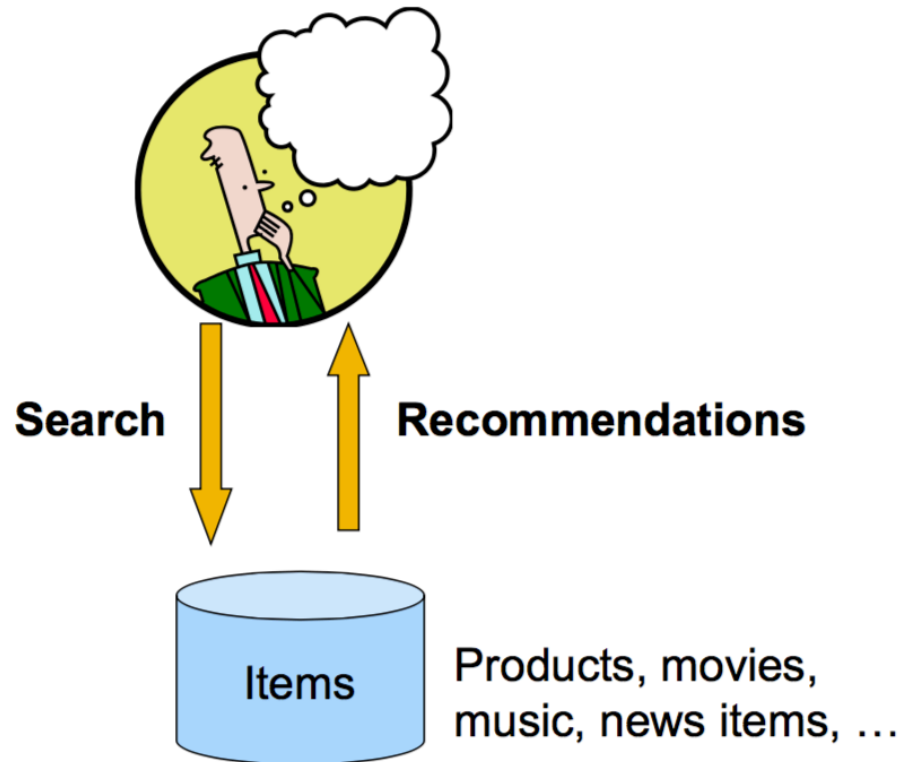
Making Recommendations: Content Filtering Collaborative Filtering Latent Factors

Jay Urbain, Ph.D.

Electrical Engineering and Computer Science Department
Milwaukee School of Engineering

Credits: See last page with references


Recommendations



From scarcity to abundance

- Shelf space is a scarce commodity for traditional retailers
 - Also: TV networks, movie theaters,...
- The web enables near zero-cost dissemination of information about products
 - From scarcity to abundance
 - Gives rise to the “Long Tail” phenomenon (Wired)

Long Tail



Long Tail

- More choice necessitates better filters
- Recommendation engines
- How **Into Thin Air** made **Touching the Void** a bestseller (<http://www.wired.com/wired/archive/12.10/tail.html>)
- Examples
 - Books, movies, music, new car sales, medical treatments
 - People (friend recommendations on Facebook, LinkedIn, and Twitter)

Types of recommendations

- Editorial and hand curated
 - List of favorites
 - Lists of “essential” items
- Simple aggregates
 - Top 10, Most Popular, Recent Uploads
- **Tailored to individual users**
 - Amazon, Netflix, Pandora ...
 - Our focus here

Recommendation Model

- **C** = set of **Customers**
- **S** = set of **Items**
- **Utility function $u: C \times S \rightarrow R$**
 - **R** = set of ratings
 - **R** is a totally ordered set, e.g., **0-5** stars, real number in **$[0,1]$**

Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Problems

- Gathering “known” ratings for matrix
 - How to collect the data in the utility matrix
- Extrapolate unknown ratings from the known ones
 - Mainly interested in high unknown ratings
 - Not typically interested in knowing what you don’t like but what you like
- Evaluating extrapolation methods
 - How to measure success/performance of recommendation methods

Gathering Ratings

- Explicit
 - Ask people to rate items
 - Doesn't scale: only a small fraction of users leave ratings and reviews
- Implicit
 - Learn ratings from user actions
 - E.g., purchase implies high rating
 - What about low ratings?

Extrapolating Utilities

- Key problem: matrix \mathbf{U} is **sparse**
- Most people have not rated most items
- Cold start:
 - New items have no ratings
 - New users have no history
- Three approaches to recommender systems
 1. Content-based
 2. Collaborative
 3. Latent factor based

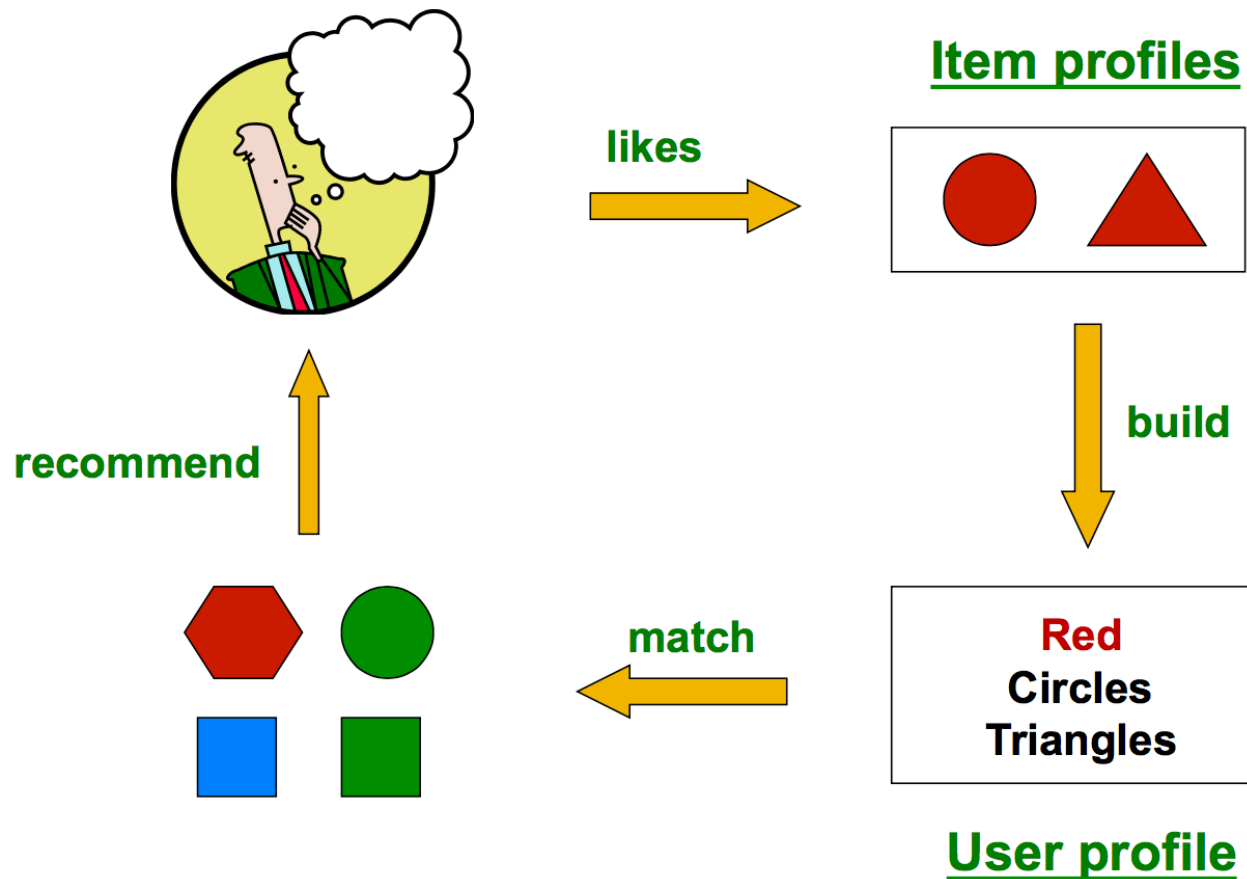
Content-based

Main idea: Recommend items to customer x similar to previous items rated highly by x

Examples:

- Movies
 - Same actor(s), director, genre, ...
- Websites, blogs, news
 - Articles with “similar” content
- People
 - Recommend people with many common friends

Build item profiles, make recommendations from profile



Item profiles

- For each item, create an **item profile**
- Profile is a set of features
 - **Movies:** author, title, actor, director,...
 - **Images, videos:** metadata and tags
 - **People:** Set of friends
- Think of the item profile as a feature vector
 - One attribute per feature per item (e.g., each actor, director,...)
 - Vector might be boolean or real-valued or mix
- Predict how similar any item is to any other item

User profiles

- User has rated items with profiles i_1, \dots, i_n
- Simple: (weighted) average of rated **item profiles**
- Variant: Normalize weights using average rating of user
- More sophisticated aggregations possible – learn feature weights.

Example 1: Boolean Utility Matrix

- Items are movies, only feature is “Actor”
 - Item profile: vector with 0 or 1 for each Actor
- Suppose user x has watched 5 movies
 - 2 movies featuring actor A
 - 3 movies featuring actor B
- User profile = mean of item profiles
 - Feature A's weight = $2/5 = 0.4$
 - Feature B's weight = $3/5 = 0.6$

Example 2: Star Ratings

- Same example, 1-5 star ratings
 - Actor A's movies rated 3 and 5
 - Actor B's movies rated 1, 2 and 4
- Useful step: Normalize ratings by subtracting user's mean rating (mean=3)
 - Actor A's normalized ratings = 0, +2
 - Profile weight = $(0 + 2)/2 = 1$
 - Actor B's normalized ratings = -2, -1, +1
 - Profile weight = $-2/3$

Making Predictions

- User profile \mathbf{x} , Item profile \mathbf{i}
- Estimate $U(\mathbf{x}, \mathbf{i}) = \cos(\theta) = (\mathbf{x} \cdot \mathbf{i}) / (|\mathbf{x}| |\mathbf{i}|)$
- Or other similarity measurement

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Pros: Content-based approach


- No need for data on other users
- Able to recommend to users with unique tastes
- Able to recommend new & unpopular items
- No first-rater problem
- Explanations for recommended items
- Content features that caused an item to be recommended

Cons: Content-based approach

- Finding the appropriate features is hard
 - E.g., images, movies, music
- Overspecialization
 - Never recommends items outside user's content profile
 - People might have multiple interests
 - Unable to exploit quality judgments of other users
- Cold-start problem for new users
 - How to build a user profile?

Collaborative Filtering

- Consider user x
- Find set N of other users whose ratings are “similar” to x ’s ratings
- Estimate x ’s ratings based on ratings of users in N



Similar Users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- Consider users \mathbf{x} and \mathbf{y} with rating vectors \mathbf{r}_x and \mathbf{r}_y
- We need a similarity metric $\text{sim}(\mathbf{x}, \mathbf{y})$
- Capture intuition that $\text{sim}(A,B) > \text{sim}(A,C)$

Jaccard Similarity

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- $\text{sim}(A,B) = |r_A \cap r_B| / |r_A \cup r_B|$
- $\text{sim}(A,B) = 1/5$; $\text{sim}(A,C) = 2/4$
 - $\text{sim}(A,B) < \text{sim}(A,C)$
- Problem: Ignores rating values!

Cosine Similarity

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

```
// sim(A,B)
```

```
>>> (5*4)/
```

```
(math.sqrt(4**2+5**2+1**2)*math.sqrt(5**2+5
```

```
**2+4**2))
```

```
0.379
```

- $\text{sim}(A,B) = \cos(r_A, r_B)$

- $\text{sim}(A,B) = 0.38, \text{sim}(A,C) = 0.32$

- $\text{sim}(A,B) < \text{sim}(A,C)$, but not by much

- Problem: treats missing ratings as negative

Centered Cosine – Pearson Correlation Coefficient

- Normalize ratings by subtracting row mean

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	$2/3$			$5/3$	$-7/3$		
<i>B</i>	$1/3$	$1/3$	$-2/3$				
<i>C</i>				$-5/3$	$1/3$	$4/3$	
<i>D</i>		0					0

Centered Cosine

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	-7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

- $\text{sim}(A,B) = \cos(r_A, r_B) = 0.09$; $\text{sim}(A,C) = -0.56$
 - $\text{sim}(A,B) > \text{sim}(A,C)$
- Captures intuition better
 - Missing ratings treated as “average”
 - Handles “tough raters” and “easy raters”
- Also known as **Pearson Correlation**

Rating Predictions

- Let r_x be the vector of user x 's ratings
- Let N be the set of k users most similar to x who have also rated item i
- Prediction for user x and item i
- Option 1: $r_{xi} = 1/k \sum_{y \in N} r_{yi}$
- Option 2: $r_{xi} = \sum_{y \in N} s_{xy} r_{yi} / \sum_{y \in N} s_{xy}$
where $s_{xy} = \text{sim}(x,y)$

Item-Item Collaborative Filtering

- So far: **User-user collaborative filtering**
- **Another view: Item-item**
 - For item i , find other similar items
 - Estimate rating for item i based on ratings for similar items
 - Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

s_{ij} ... similarity of items i and j
 r_{xj} ... rating of user x on item j
 $N(i;x)$... set items rated by x similar to i

Item-Item CF (n=2)

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3			5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- unknown rating



- rating between 1 to 5

Item-Item CF (n=2)

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

Item-Item CF (n=2)

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	sim(1,m)
movies	1	1		3		?	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		<u>0.59</u>

Neighbor selection:
Identify movies similar to
movie 1, **rated by user 5**

Here we use Pearson correlation as similarity:

1) Subtract mean rating m_i from each movie i

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]

2) Compute cosine similarities between rows

Item-Item CF (n=2)

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	
movies	1	1		3		?	5			5		4		sim(1,m) 1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		<u>0.59</u>

Compute similarity weights:

$s_{13}=0.41$, $s_{16}=0.59$

Item-Item CF (n=2)

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		2.6	5			5		4	
	2			5	4			4			2	1	3
	<u>3</u>	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	<u>6</u>	1		3		3			2			4	

Predict by taking weighted average:

$$r_{15} = (0.41*2 + 0.59*3) / (0.41+0.59) = 2.6$$

Item-item vs. User-user

- In theory, user-user and item-item are dual approaches
- In practice, item-item outperforms user-user in many use cases
- Items are “simpler” than users
 - Items belong to a small set of “genres,” users have varied tastes
 - Item Similarity is more meaningful than User Similarity